

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings('ignore' )
import plotly.express as px
```

```
In [2]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc
from sklearn.model_selection import GridSearchCV
from imblearn.over_sampling import SMOTE
import scikitplot as skplt
```

```
In [3]: from sklearn.svm import SVC
from xgboost import XGBClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier, AdaBoostClassifier
from sklearn.preprocessing import LabelEncoder, MinMaxScaler
from sklearn.model_selection import train_test_split
from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import RandomOverSampler
```

```
In [4]: sns.set_style('darkgrid')
plt.rcParams['figure.figsize']=(15,8)
plt.rcParams['font.size']=18
```

```
In [5]: # importing data frame
df = pd.read_csv("stroke.csv")
```

```
In [6]: # top 5 row of df
df.head()
```

```
Out[6]:
```

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level
0	9046	Male	67.0	0	1	Yes	Private	Urban	228.69
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	202.21
2	31112	Male	80.0	0	1	Yes	Private	Rural	105.92
3	60182	Female	49.0	0	0	Yes	Private	Urban	171.23
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural	174.12

```
In [7]: # shape of dataframe
df.shape
```

```
Out[7]: (5110, 12)
```

```
In [8]: # columns of df
df.columns
```

```
Out[8]: Index(['id', 'gender', 'age', 'hypertension', 'heart_disease', 'ever_married',
        'work_type', 'Residence_type', 'avg_glucose_level', 'bmi',
```

```
'smoking_status', 'stroke'],
dtype='object')
```

```
In [9]: # basic information
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    5110 non-null   int64
1   gender                5110 non-null   object
2   age                   5110 non-null   float64
3   hypertension          5110 non-null   int64
4   heart_disease         5110 non-null   int64
5   ever_married          5110 non-null   object
6   work_type             5110 non-null   object
7   Residence_type        5110 non-null   object
8   avg_glucose_level     5110 non-null   float64
9   bmi                   4909 non-null   float64
10  smoking_status        5110 non-null   object
11  stroke                5110 non-null   int64
dtypes: float64(3), int64(4), object(5)
memory usage: 479.2+ KB
```

```
In [10]: #duplicated rows
df.duplicated().sum()
```

```
Out[10]: 0
```

observation

1. No duplicated row present in df

```
In [11]: # null values:
df.isnull().sum()
```

```
Out[11]: id                    0
gender                0
age                   0
hypertension          0
heart_disease         0
ever_married          0
work_type             0
Residence_type        0
avg_glucose_level     0
bmi                   201
smoking_status        0
stroke                0
dtype: int64
```

1. only BMI columns has null values

```
In [12]: #description of df
df.describe()
```

```
Out[12]:
```

	id	age	hypertension	heart_disease	avg_glucose_level	bmi	stroke
count	5110.000000	5110.000000	5110.000000	5110.000000	5110.000000	4909.000000	5110.000000
mean	36517.829354	43.226614	0.097456	0.054012	106.147677	28.893237	0.048728
std	21161.721625	22.612647	0.296607	0.226063	45.283560	7.854067	0.215320

min	67.000000	0.080000	0.000000	0.000000	55.120000	10.300000	0.000000
25%	17741.250000	25.000000	0.000000	0.000000	77.245000	23.500000	0.000000
50%	36932.000000	45.000000	0.000000	0.000000	91.885000	28.100000	0.000000
75%	54682.000000	61.000000	0.000000	0.000000	114.090000	33.100000	0.000000
max	72940.000000	82.000000	1.000000	1.000000	271.740000	97.600000	1.000000

```
In [13]: # dropping unwanted columns:
df.drop("id",1,inplace=True)
```

```
In [14]: # no of unique values:
df.nunique()
```

```
Out[14]: gender                3
age                104
hypertension        2
heart_disease        2
ever_married        2
work_type            5
Residence_type      2
avg_glucose_level  3979
bmi                 418
smoking_status       4
stroke              2
dtype: int64
```

```
In [15]: # printing unique values:
for i in df.columns:
    print(df[i].value_counts())
    print("-----")
```

```
Female    2994
Male      2115
Other      1
Name: gender, dtype: int64
```

```
-----
78.00     102
57.00      95
52.00      90
54.00      87
51.00      86
```

```
...
1.40       3
0.48       3
0.16       3
0.40       2
0.08       2
```

```
Name: age, Length: 104, dtype: int64
```

```
-----
0      4612
1       498
```

```
Name: hypertension, dtype: int64
```

```
-----
0      4834
1       276
```

```
Name: heart_disease, dtype: int64
```

```
-----
Yes      3353
No       1757
```

```
Name: ever_married, dtype: int64
```

```
-----
Private      2925
```

```

Self-employed      819
children           687
Govt_job           657
Never_worked       22
Name: work_type, dtype: int64
-----
Urban      2596
Rural      2514
Name: Residence_type, dtype: int64
-----
93.88      6
91.68      5
91.85      5
83.16      5
73.00      5
..
111.93     1
94.40      1
95.57      1
66.29      1
85.28      1
Name: avg_glucose_level, Length: 3979, dtype: int64
-----
28.7      41
28.4      38
26.7      37
27.6      37
26.1      37
..
48.7      1
49.2      1
51.0      1
49.4      1
14.9      1
Name: bmi, Length: 418, dtype: int64
-----
never smoked      1892
Unknown           1544
formerly smoked    885
smokes            789
Name: smoking_status, dtype: int64
-----
0      4861
1       249
Name: stroke, dtype: int64
-----

```

```

In [16]: # numerical columns
num_cols = df.select_dtypes(include=['int','float']).columns
# categorical columns:
cat_cols = df.select_dtypes(include=['object']).columns

```

```

In [17]: num_cols

```

```

Out[17]: Index(['age', 'hypertension', 'heart_disease', 'avg_glucose_level', 'bmi',
               'stroke'],
              dtype='object')

```

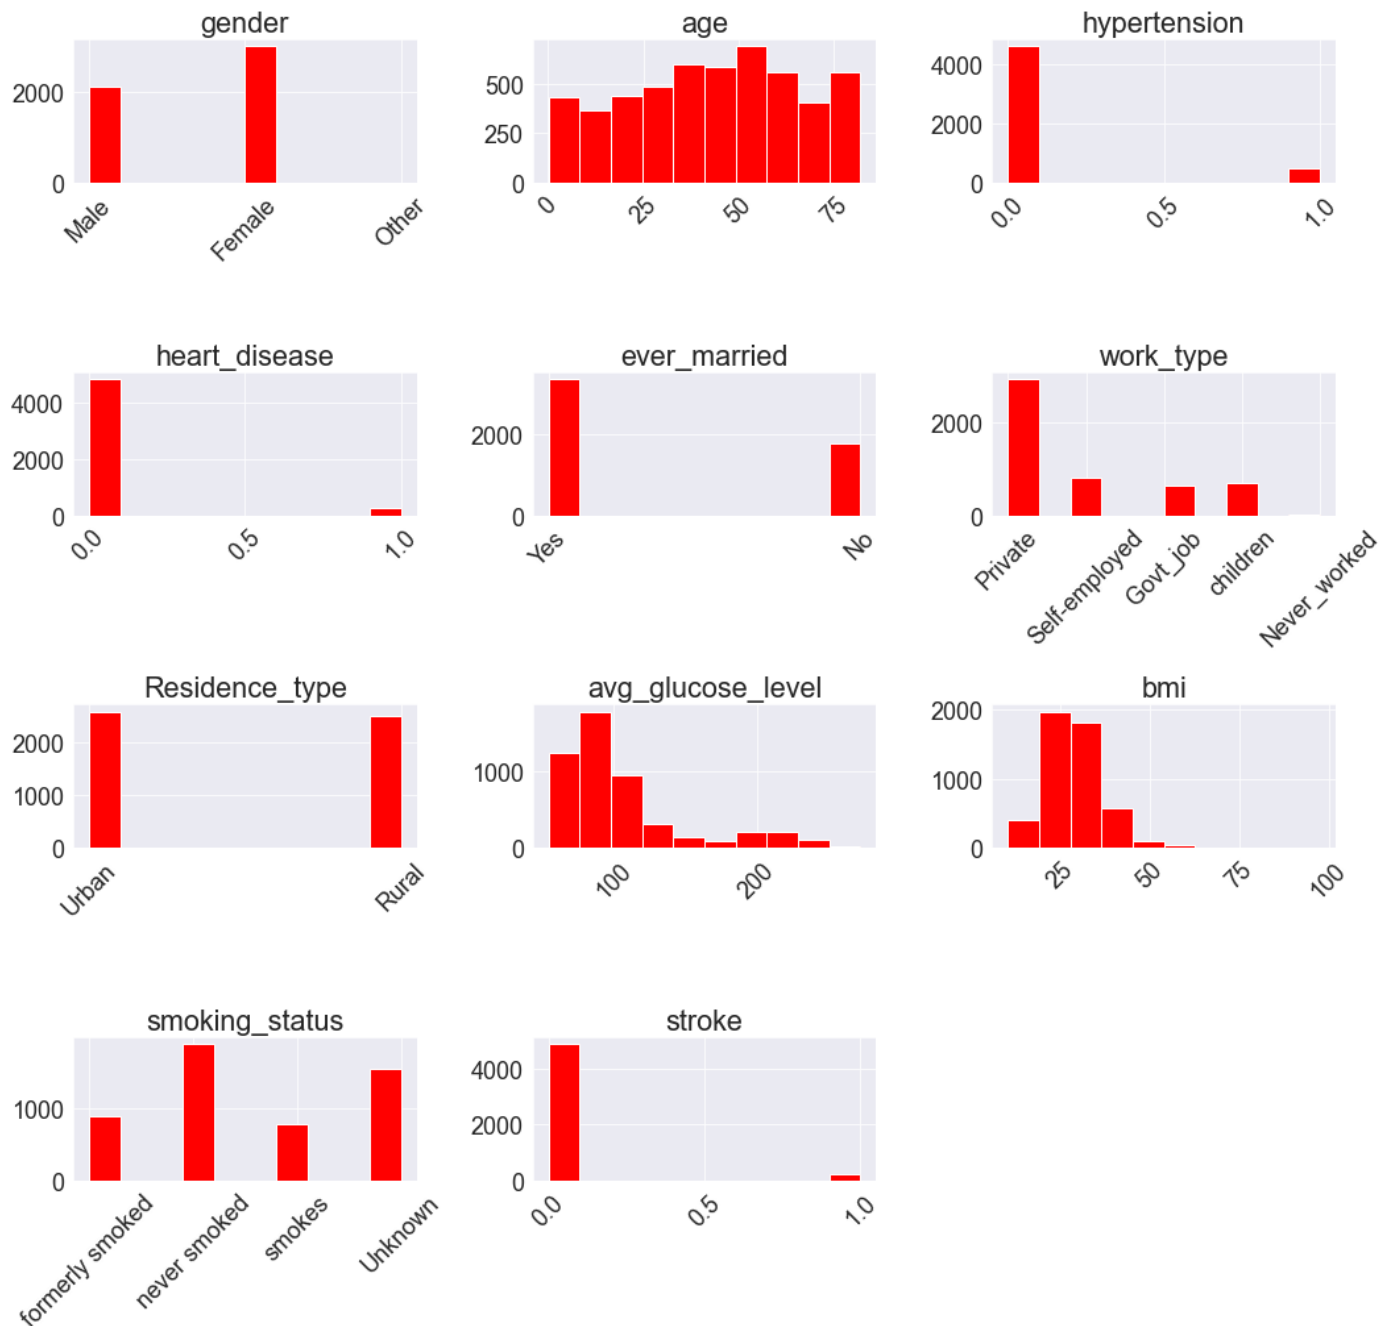
```

In [18]: plt.figure(figsize=(15,15))
plt.suptitle('Univariate Analysis using histplot', fontsize=20, fontweight='bold', alpha
for j,i in enumerate(df):
    plt.subplot(4,3,j+1)
    plt.title(i)
    plt.xticks(rotation=45)
    plt.hist(df[i],color='red')

```

```
plt.tight_layout()
plt.show()
```

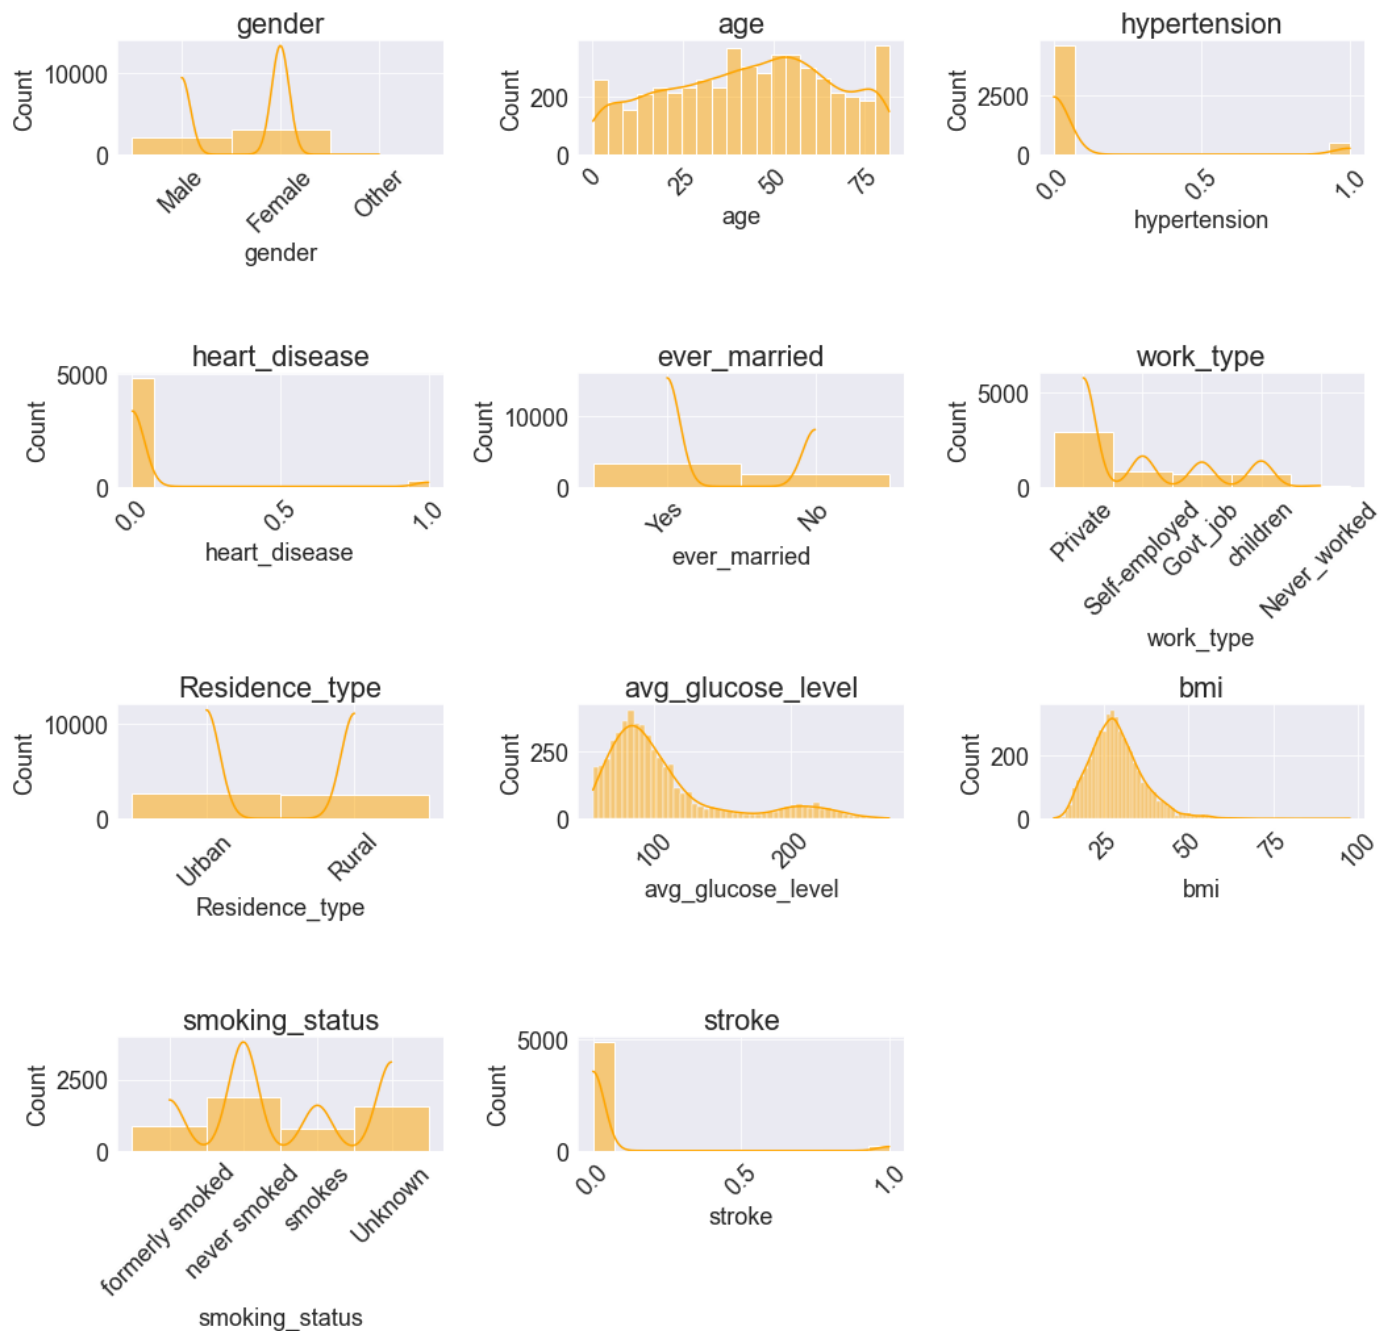
Univariate Analysis using histplot



```
In [19]: plt.figure(figsize=(15,15) )
plt.suptitle('Univariate Analysis', fontsize=20, fontweight='bold', alpha=0.8, y=1.)
for j,i in enumerate(df):
    plt.subplot(4,3,j+1)
    plt.title(i)
    plt.xticks(rotation=45)
    sns.histplot(df[i],color='orange',kde=True)

plt.tight_layout()
plt.show()
```

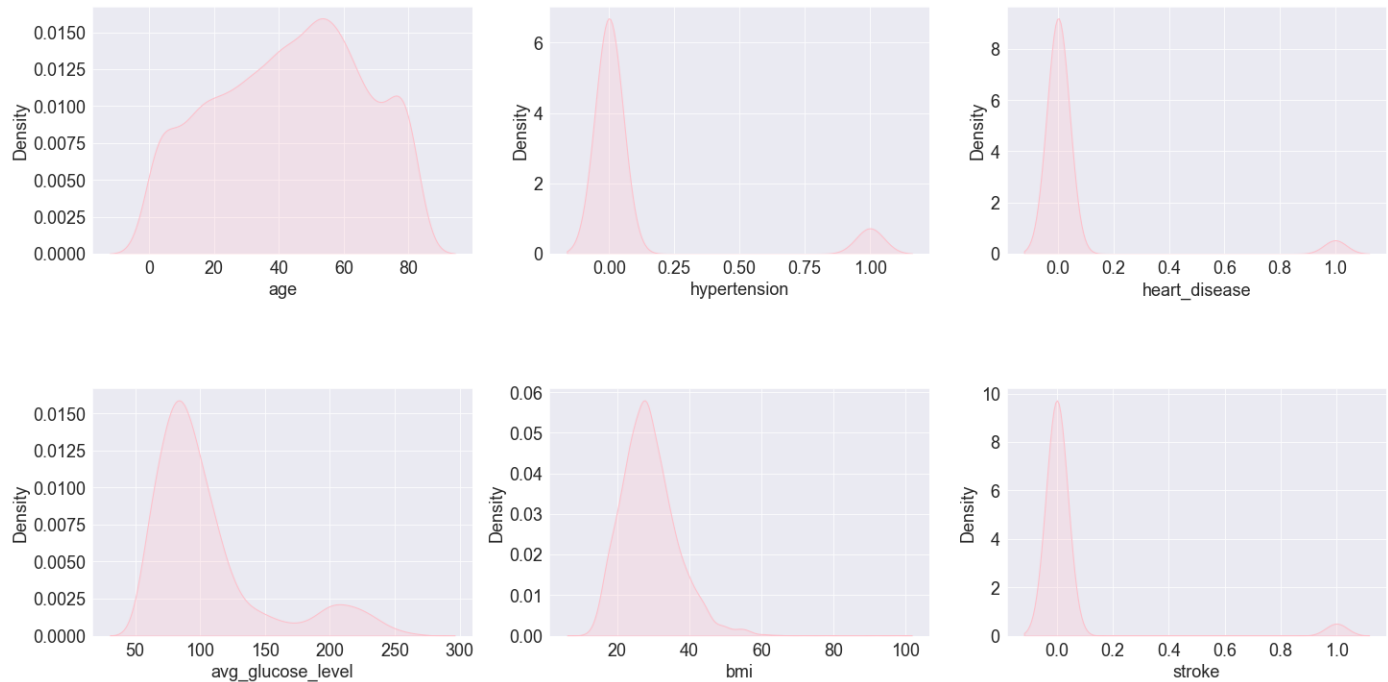
Univariate Analysis



```
In [20]: plt.suptitle('Univariate Analysis of Numerical Features using KDE Plot', fontsize=20, fo
fig, ax = plt.subplots(ncols=3, nrows=2, figsize=(20,10))
index = 0
ax = ax.flatten()

for col in num_cols:
    sns.kdeplot(x=col, data=df, ax=ax[index], shade=True, color="Pink")
    index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```

<Figure size 1080x576 with 0 Axes>



In []:

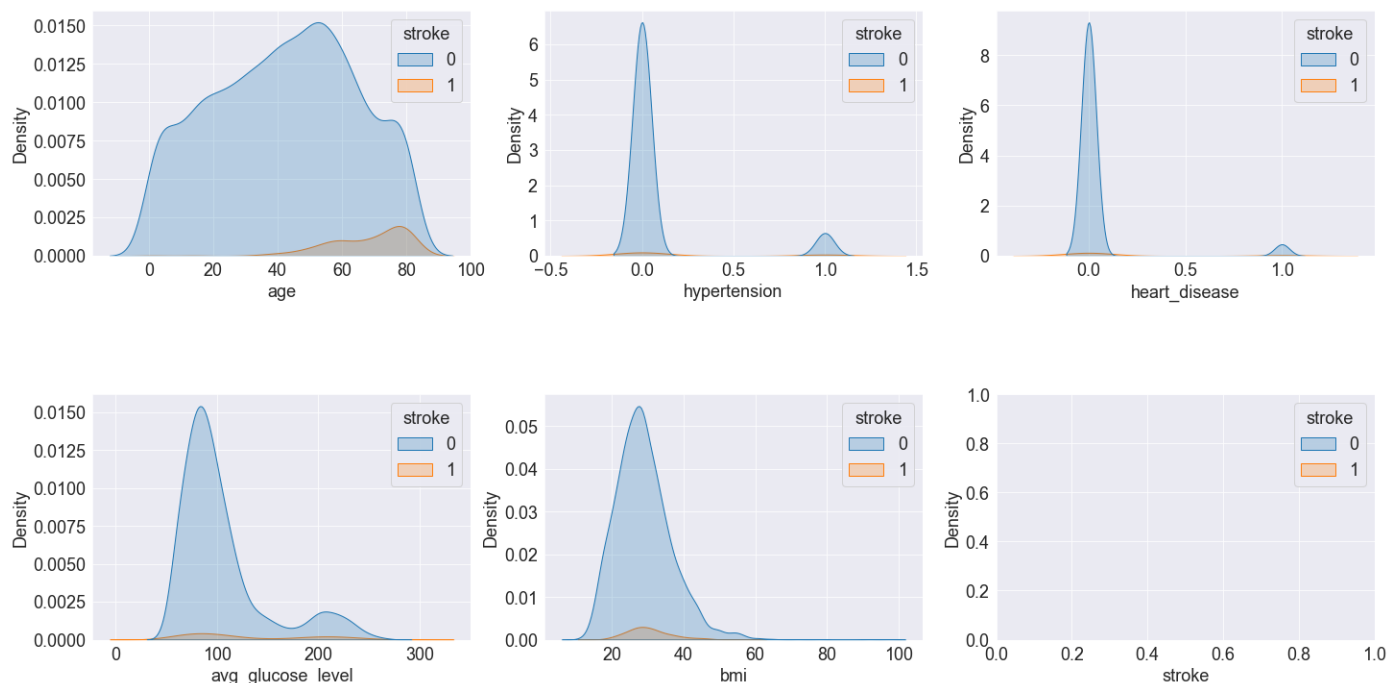
observation

1. hypertension , heart_disease and stroke are left skewed
2. avg_glucose_level and BMI is normal distributed

```
In [21]: plt.suptitle('Univariate Analysis of Numerical Features using KDE Plot', fontsize=20, fo
fig, ax = plt.subplots(ncols=3, nrows=2, figsize=(20,10))
index = 0
ax = ax.flatten()

for col in num_cols:
    sns.kdeplot(x=col, data=df, ax=ax[index], shade=True, color="Pink", hue='stroke')
    index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```

<Figure size 1080x576 with 0 Axes>



```
In [22]: plt.suptitle('Univariate Analysis of categorical Features using count plot', fontsize=20)
```

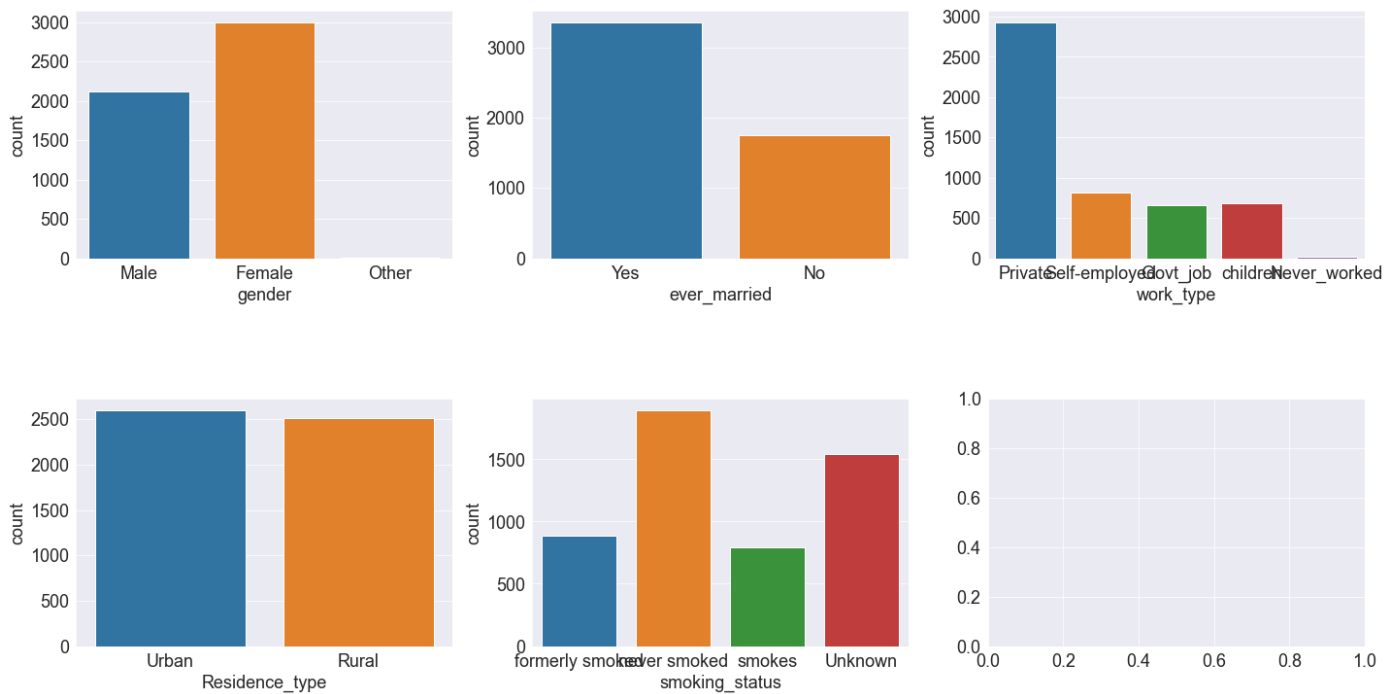
```

fig, ax = plt.subplots(ncols=3, nrows=2, figsize=(20,10))
index = 0
ax = ax.flatten()

for col in cat_cols:
    sns.countplot(x=col, data=df, ax=ax[index])
    index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)

```

<Figure size 1080x576 with 0 Axes>

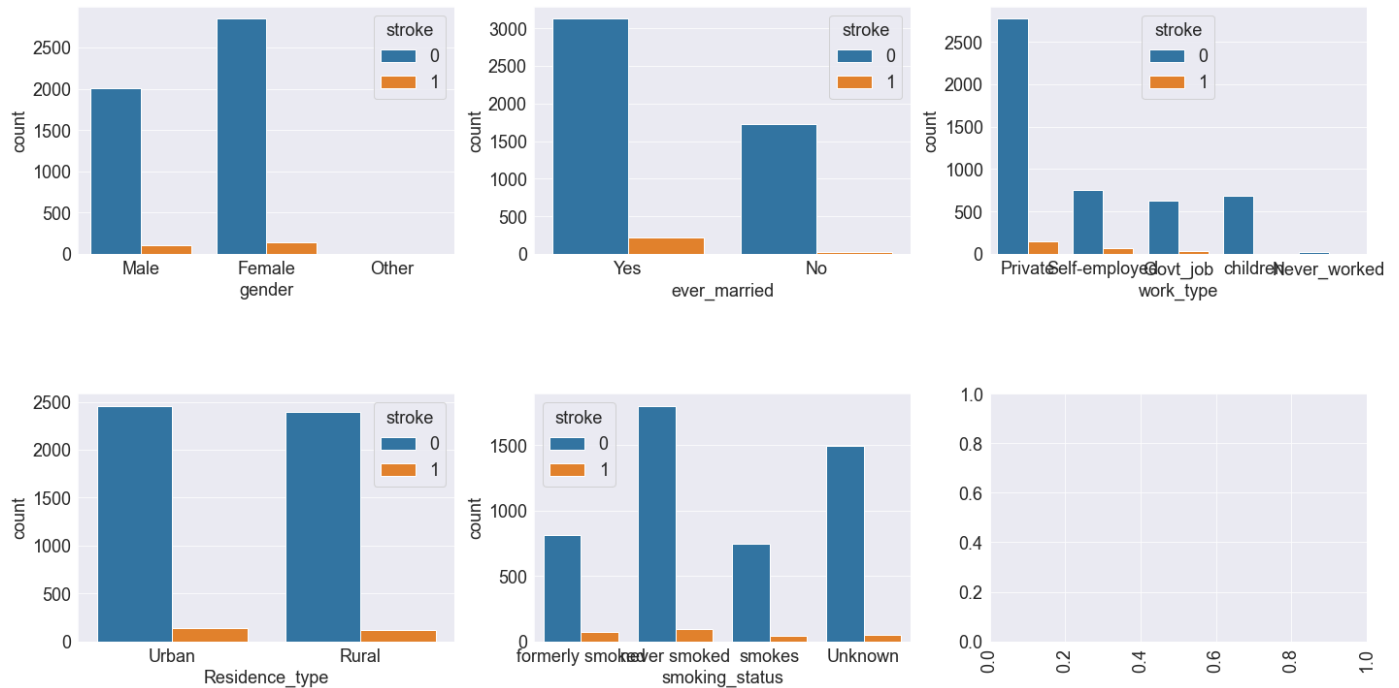


```

In [23]: plt.suptitle('Univariate Analysis of categorical Features using count plot', fontsize=20)
fig, ax = plt.subplots(ncols=3, nrows=2, figsize=(20,10))
index = 0
ax = ax.flatten()
for col in cat_cols:
    sns.countplot(x=col, data=df, ax=ax[index], hue='stroke')
    index += 1
plt.xticks(rotation=90)
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
plt.show()

```

<Figure size 1080x576 with 0 Axes>



1. gender has 3 unique values male,female,and other

```
In [24]: df['gender'].unique()
```

```
Out[24]: array(['Male', 'Female', 'Other'], dtype=object)
```

```
In [25]: df['gender'].value_counts()
```

```
Out[25]: Female      2994
Male        2115
Other         1
Name: gender, dtype: int64
```

```
In [26]: df[df['gender']=='Other']
```

```
Out[26]:
```

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi
3116	Other	26.0	0	0	No	Private	Rural	143.33	22.1

```
In [27]: df = df.drop(df[df['gender']=='Other'].index)
```

```
In [28]: df[df['gender']=='Other']
```

```
Out[28]:
```

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	sn
--	--------	-----	--------------	---------------	--------------	-----------	----------------	-------------------	-----	----

```
In [29]: # handling missing values
print('Mean of BMI = ',df['bmi'].mean())
print("Median of BMI= ",df['bmi'].median())
```

```
Mean of BMI = 28.894559902200502
Median of BMI= 28.1
```

```
In [30]: # filling missing values with mean
bmi_mean=df['bmi'].mean()
df['bmi']=df['bmi'].fillna(bmi_mean)
```

```
In [31]: df['bmi'].isnull().sum()
```


Out[31]: 0

```
In [32]: # counting the number of passengers who are satisfied and who Dissatisfied
pie_df=pd.DataFrame(df.groupby('stroke')['stroke'].count())
pie_df
```

Out[32]:


stroke	
0	4860
1	249

```
In [33]: colors = ['#B2BEBF', '#BD2A2E']
plt.pie(pie_df['stroke'], labels=['no stroke', 'stroke'],
        autopct='%.1f%%', colors=colors, radius=2, explode = (0, 0.1), shadow=True)
plt.show()
```




```
In [34]: sns.set_theme(style="darkgrid")
fig = plt.figure(figsize=(9,7))
sns.histplot(df['avg_glucose_level'], color="#2C3532", label="avg_glucose_level", kde=True)
plt.legend()
```

Out[34]: <matplotlib.legend.Legend at 0x28738359700>




```
In [35]: fig = plt.figure(figsize=(9,7))
sns.histplot(df['bmi'], color="#0F6466", label="bmi", kde= True)
plt.legend()
plt.show()
```




```
In [36]: plt.figure(figsize=(12,7))
sns.histplot(
    df, x="bmi", hue="stroke",
    element="step",
    stat="density", common_norm=False,
```

```
)
plt.title('Stroke vs no stroke by BMI', fontsize=15)
plt.show()
```

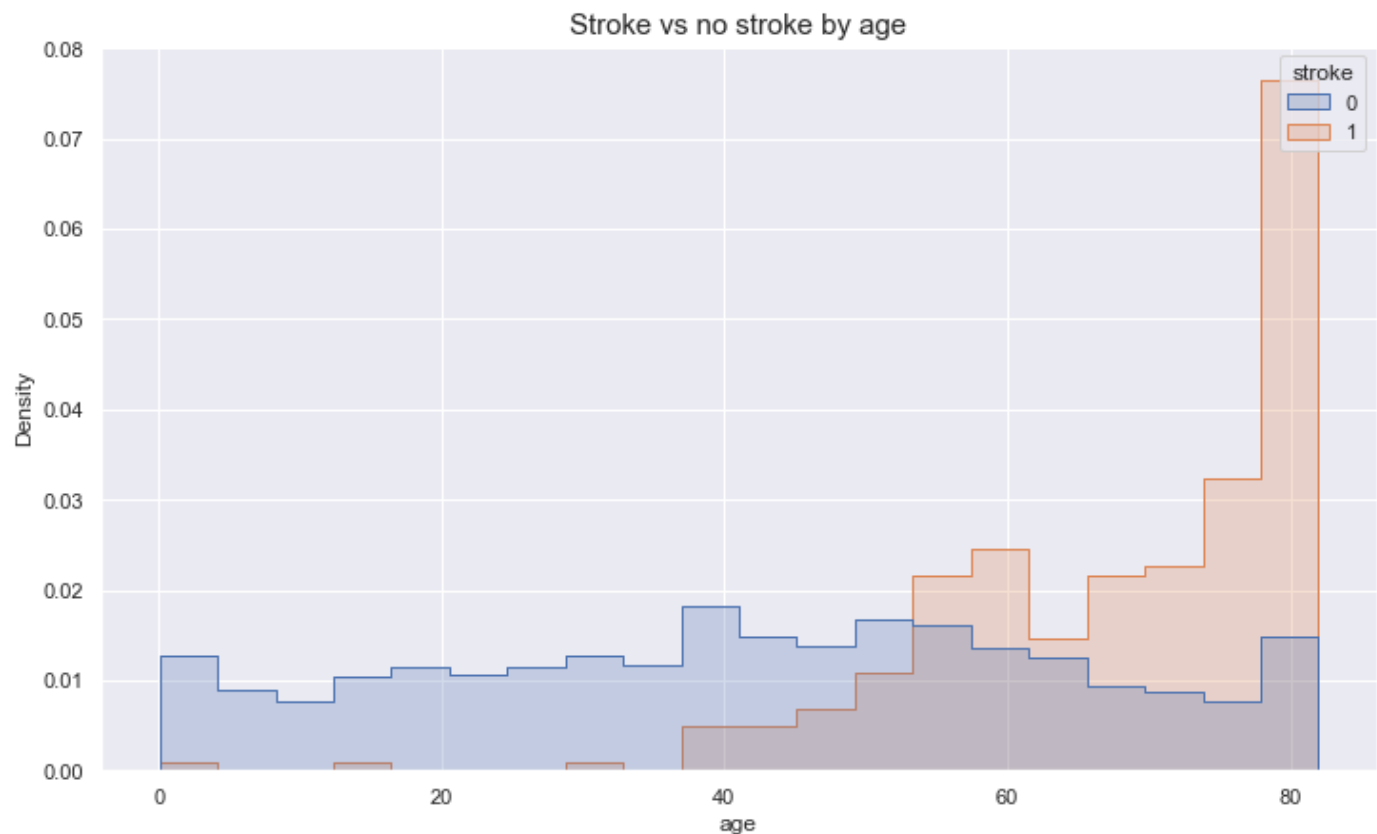


```
In [37]: plt.figure(figsize=(12,7))
sns.histplot(
    df, x="avg_glucose_level", hue="stroke",
    element="step",
    stat="density", common_norm=False,
)
plt.title('Stroke vs no stroke by avg_glucose_level', fontsize=15)
plt.show()
```

Stroke vs no stroke by avg_glucose_level



```
In [38]: plt.figure(figsize=(12,7))
sns.histplot(
    df, x="age", hue="stroke",
    element="step",
    stat="density", common_norm=False,
)
plt.title('Stroke vs no stroke by age', fontsize=15)
plt.show()
```



```
In [39]: plt.suptitle('Univariate Analysis of Numerical Features using KDE Plot', fontsize=20, fo
fig, ax = plt.subplots(ncols=3, nrows=2, figsize=(20,10))
```

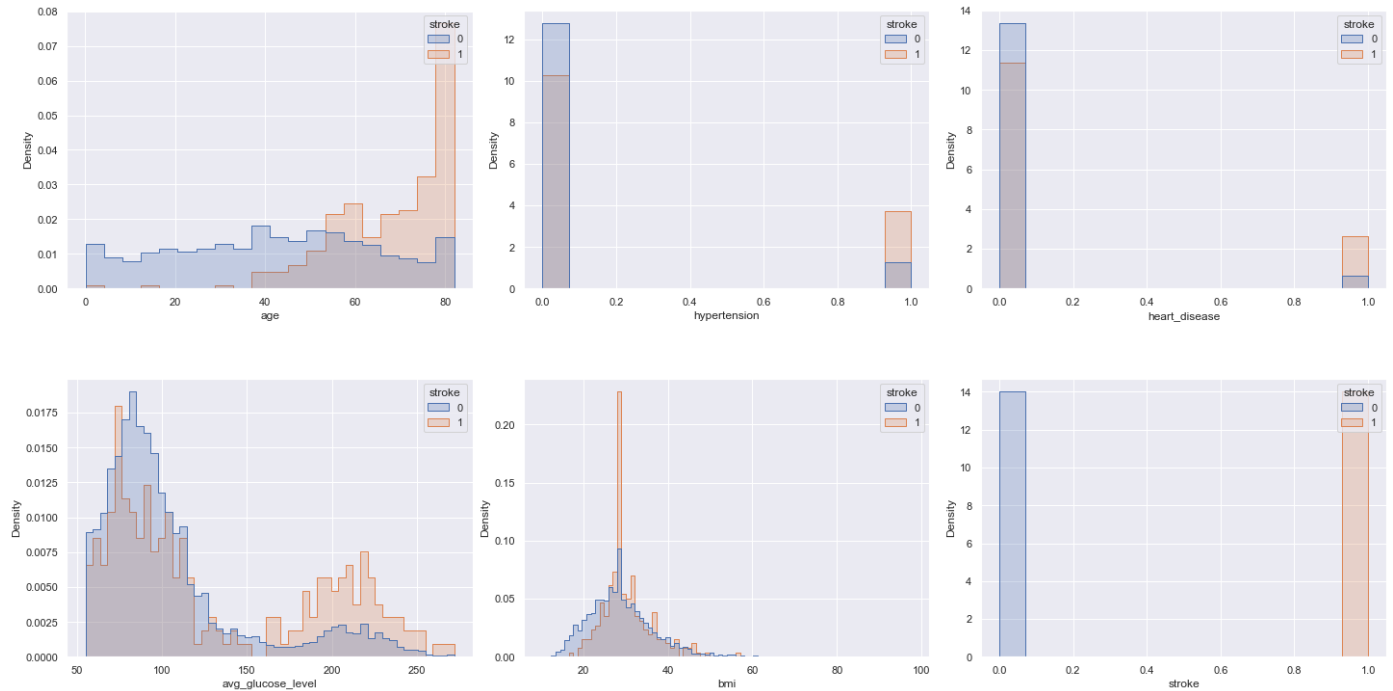
```

index = 0
ax = ax.flatten()

for col in num_cols:
    sns.histplot(x=col, data=df, ax=ax[index], hue='stroke', element='step', stat='dens')
    index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)

```

<Figure size 1080x576 with 0 Axes>



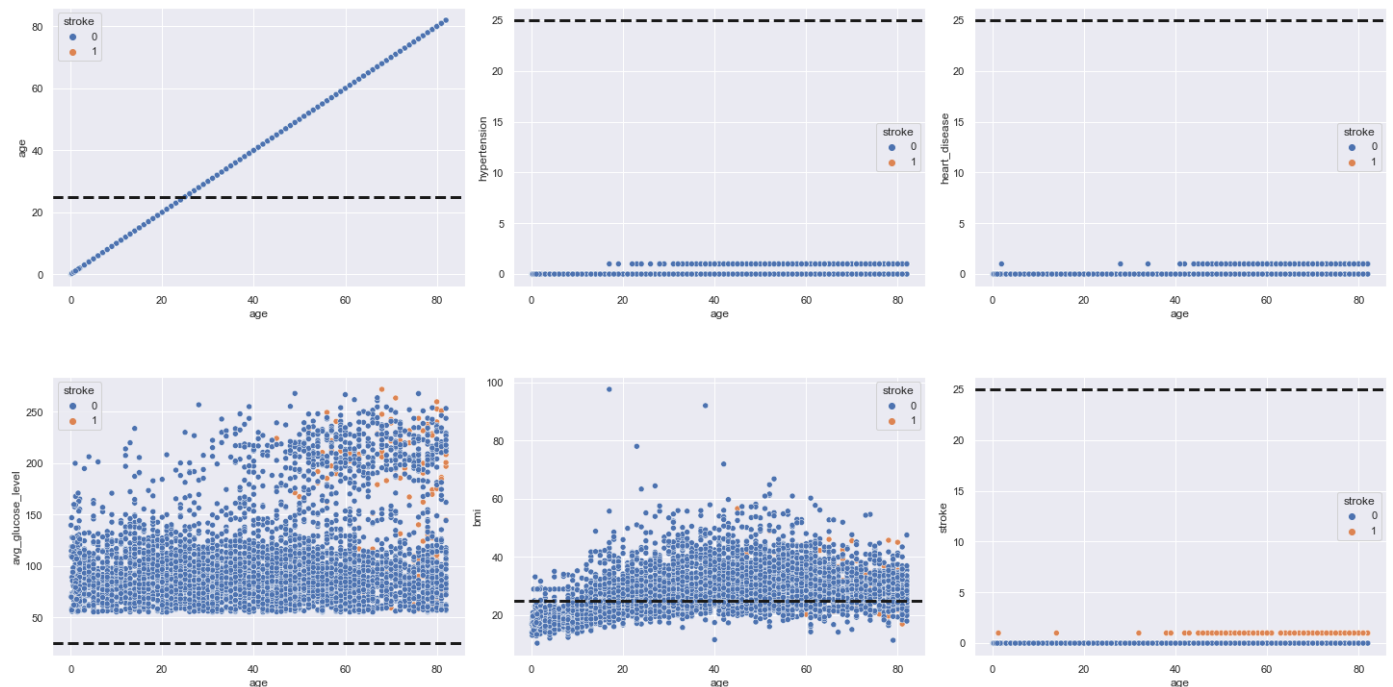
```

In [40]: plt.suptitle('Univariate Analysis of Numerical Features using KDE Plot', fontsize=20, fo
fig, ax = plt.subplots(ncols=3, nrows=2, figsize=(20,10))
index = 0
ax = ax.flatten()

for col in num_cols:
    sns.scatterplot(x='age', y = col, data=df, ax=ax[index], hue='stroke').axhline(y=
    index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)

```

<Figure size 1080x576 with 0 Axes>



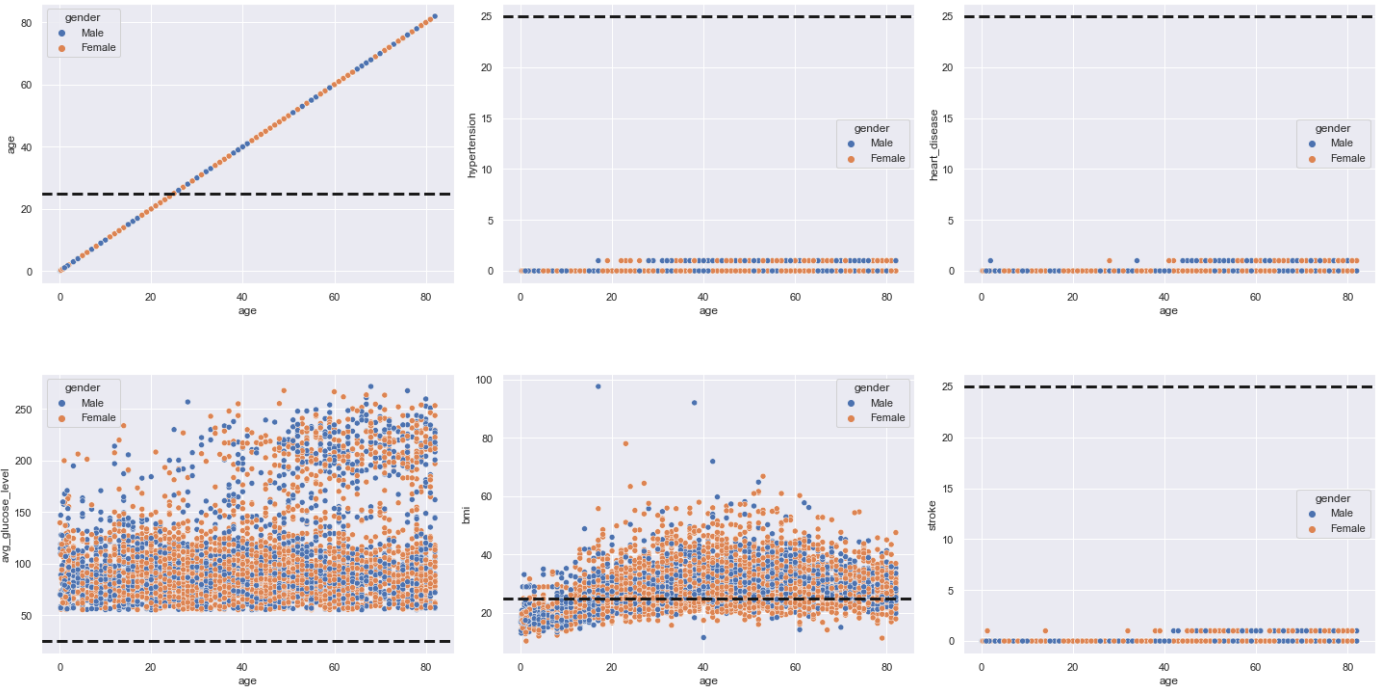
```

In [41]: plt.suptitle('Univariate Analysis of Numerical Features using KDE Plot', fontsize=20, fo
fig, ax = plt.subplots(ncols=3, nrows=2, figsize=(20,10))
index = 0
ax = ax.flatten()

for col in num_cols:
    sns.scatterplot(x='age', y = col, data=df, ax=ax[index], hue='gender').axhline(y=
    index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
plt.show()

```

<Figure size 1080x576 with 0 Axes>



```

In [42]: plt.figure(figsize=(8,6))
fig = sns.scatterplot(data=df, x="age", y="bmi", hue='gender')
fig.axhline(y= 25, linewidth=3, color='k', linestyle='--')
plt.show()

```




```

In [43]: plt.figure(figsize=(8,6))
fig = sns.scatterplot(data=df, x="age", y="avg_glucose_level", hue='gender')

```

```
fig.axhline(y=150, linewidth=3, color='k', linestyle='--')
plt.show()
```



```
In [44]: # label encoding
le = LabelEncoder()
for i in cat_cols:
    df[i] = le.fit_transform(df[i])
```

```
In [45]: df.head()
```

```
Out[45]:
```

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bn
0	1	67.0	0	1	1	2	1	228.69	36.6000
1	0	61.0	0	0	1	3	0	202.21	28.8945
2	1	80.0	0	1	1	2	0	105.92	32.5000
3	0	49.0	0	0	1	2	1	171.23	34.4000
4	0	79.0	1	0	1	3	0	174.12	24.0000

```
In [46]: # saving cleaning df
df.to_csv("clean_heart_stroke.csv")
```

```
In [47]: x = df.drop("stroke" , axis = 1).values
y = df["stroke"]
```

```
In [48]: x_train, x_test, y_train , y_test = train_test_split(x,y, test_size=0.2 , random_state=4
```

```
In [49]: sc=MinMaxScaler()
x_train = sc.fit_transform(x_train)
```


```
In [50]: # show the value counts of the calsses in the target
#we can find data impalance
print(df['stroke'].value_counts())
df['stroke'].value_counts().sort_index().plot.bar()
```

```
0    4860
```



```
1      249
Name: stroke, dtype: int64
<AxesSubplot:>
```

Out[50]:



```
In [51]: # Apply oversampling
oversample = SMOTE()
x_data_balanced, y_data_balanced = oversample.fit_resample(x_train, y_train.ravel())
```

Modeling:

Logistic regression:

```
In [52]: lr = LogisticRegression()
lr.fit(x_data_balanced, y_data_balanced)
```

```
Out[52]: ▼ LogisticRegression
LogisticRegression()
```

```
In [53]: y_pred_train_lr = lr.predict(x_data_balanced)
acc_train_lr = accuracy_score(y_data_balanced, y_pred_train_lr)

y_pred_test_lr = lr.predict(x_test)
acc_test_lr = accuracy_score(y_test, y_pred_test_lr)
print(acc_train_lr)
print(acc_test_lr)

0.7901282051282051
0.060665362035225046
```

```
In [54]: print(classification_report(y_test, y_pred_test_lr))
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	960
1	0.06	1.00	0.11	62

accuracy			0.06	1022
macro avg	0.03	0.50	0.06	1022
weighted avg	0.00	0.06	0.01	1022

```
In [55]: lr_perc_score = precision_score(y_test, y_pred_test_lr)
lr_rec_score= recall_score(y_test, y_pred_test_lr)
lr_f1_score = f1_score(y_test, y_pred_test_lr)

print('Precision: %.3f' %lr_perc_score )
print('Recall: %.3f' % lr_rec_score)
print('F-measure: %.3f' % lr_f1_score)
```


```
Precision: 0.061
Recall: 1.000
F-measure: 0.114
```

```
In [56]: y_pred_prob_lr = lr.predict_proba(x_test)[: , 1]
lr_roc_auc_score = roc_auc_score(y_test, y_pred_prob_lr)
print('ROC AUC Score:', lr_roc_auc_score)
```

```
ROC AUC Score: 0.5
```

```
In [57]: skplt.metrics.plot_confusion_matrix(y_test, y_pred_test_lr, figsize=(6,6), cmap= 'Blues')
```

```
Out[57]: <AxesSubplot:title={'center':'Confusion Matrix'}, xlabel='Predicted label', ylabel='True label'>
```



Decision Tree

```
In [58]: dt =DecisionTreeClassifier(max_features=14 , max_depth=12, criterion= 'gini')
dt.fit(x_data_balanced, y_data_balanced)
```

```
Out[58]: ▼ DecisionTreeClassifier
DecisionTreeClassifier(max_depth=12, max_features=14)
```

```
In [59]: y_pred_train_dt = dt.predict(x_data_balanced)
acc_train_dt = accuracy_score(y_data_balanced, y_pred_train_dt)
```

```

y_pred_test_dt = dt.predict(x_test)
acc_test_dt = accuracy_score(y_test, y_pred_test_dt)
print(acc_train_dt)
print(acc_test_dt)

```

```

0.9401282051282052
0.901174168297456

```

```

In [60]: print(classification_report(y_test, y_pred_test_dt))

```

	precision	recall	f1-score	support
0	0.94	0.96	0.95	960
1	0.05	0.03	0.04	62
accuracy			0.90	1022
macro avg	0.49	0.49	0.49	1022
weighted avg	0.88	0.90	0.89	1022

```

In [61]: dt_perc_score = precision_score(y_test, y_pred_test_dt)
dt_rec_score= recall_score(y_test, y_pred_test_dt)
dt_f1_score = f1_score(y_test, y_pred_test_dt)

```

```

print('Precision: %.3f' % dt_perc_score)
print('Recall: %.3f' % dt_rec_score)
print('F-measure: %.3f' % dt_f1_score)

```

```

Precision: 0.047
Recall: 0.032
F-measure: 0.038

```

```

In [62]: y_pred_prob_dt = dt.predict_proba(x_test)[: , 1]
dt_roc_auc_score = roc_auc_score(y_test, y_pred_prob_dt)
print('ROC AUC Score:', dt_roc_auc_score)

```

```

ROC AUC Score: 0.4947748655913978

```

```


In [63]: skplt.metrics.plot_confusion_matrix(y_test, y_pred_test_dt, figsize=(6,6), cmap= 'YlGnBu')

```

```

Out[63]: <AxesSubplot:title={'center':'Confusion Matrix'}, xlabel='Predicted label', ylabel='True label'>

```



KNN

```
In [64]: knn = KNeighborsClassifier()
knn.fit(x_data_balanced, y_data_balanced)
```

```
Out[64]: ▼ KNeighborsClassifier
KNeighborsClassifier()
```

```
In [65]: y_pred_train_knn = knn.predict(x_data_balanced)
acc_train_knn = accuracy_score(y_data_balanced, y_pred_train_knn)

y_pred_test_knn = knn.predict(x_test)
acc_test_knn = accuracy_score(y_test, y_pred_test_knn)
print(acc_train_knn)
print(acc_test_knn)

0.9397435897435897
0.7906066536203522
```

```
In [66]: print(classification_report(y_test, y_pred_test_knn))
```

	precision	recall	f1-score	support
0	0.95	0.82	0.88	960
1	0.09	0.27	0.14	62
accuracy			0.79	1022
macro avg	0.52	0.55	0.51	1022
weighted avg	0.89	0.79	0.84	1022

```
In [67]: knn_perc_score = precision_score(y_test, y_pred_test_knn)
knn_rec_score = recall_score(y_test, y_pred_test_knn)
knn_f1_score = f1_score(y_test, y_pred_test_knn)

print('Precision: %.3f' % knn_perc_score)
print('Recall: %.3f' % knn_rec_score)
print('F-measure: %.3f' % knn_f1_score)
```

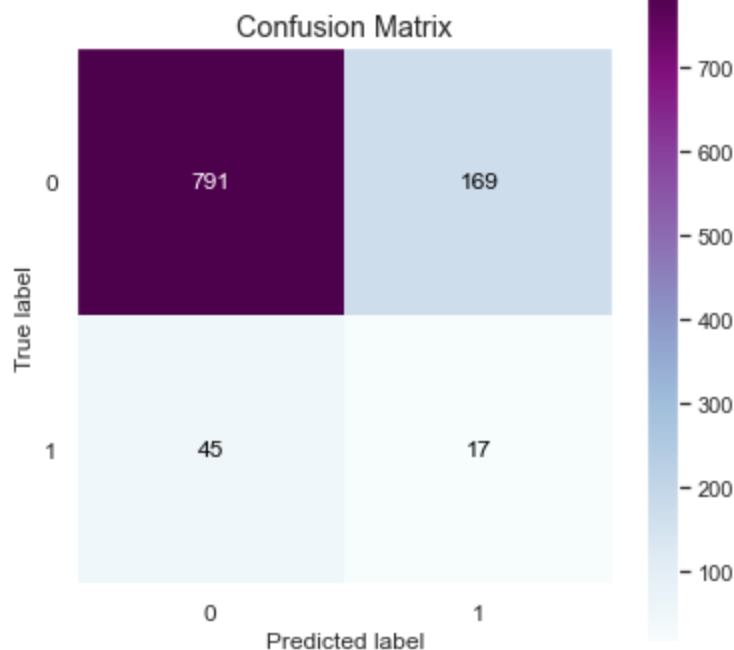
```
Precision: 0.091
Recall: 0.274
F-measure: 0.137
```

```
In [68]: y_pred_prob_knn = knn.predict_proba(x_test)[: , 1]
knn_roc_auc_score = roc_auc_score(y_test, y_pred_prob_knn)
print('ROC AUC Score:', knn_roc_auc_score)
```

```
ROC AUC Score: 0.5656418010752688
```

```
In [69]: skplt.metrics.plot_confusion_matrix(y_test, y_pred_test_knn, figsize=(6,6), cmap= 'BuPu')
```

```
Out[69]: <AxesSubplot:title={'center':'Confusion Matrix'}, xlabel='Predicted label', ylabel='True
label'>
```



SVC

```
In [70]: svc = SVC(C=100, gamma=1000, probability=True)
svc.fit(x_data_balanced, y_data_balanced)
```

```
Out[70]: SVC
SVC(C=100, gamma=1000, probability=True)
```

```
In [71]: y_pred_train_svc = svc.predict(x_data_balanced)
acc_train_svc = accuracy_score(y_data_balanced, y_pred_train_svc)

y_pred_test_svc = svc.predict(x_test)
acc_test_svc = accuracy_score(y_test, y_pred_test_svc)
print(acc_train_svc)
print(acc_test_svc)

1.0
0.9393346379647749
```

```
In [72]: print(classification_report(y_test, y_pred_test_svc))
```

	precision	recall	f1-score	support
0	0.94	1.00	0.97	960
1	0.00	0.00	0.00	62
accuracy			0.94	1022
macro avg	0.47	0.50	0.48	1022
weighted avg	0.88	0.94	0.91	1022

```
In [73]: svc_perc_score = precision_score(y_test, y_pred_test_svc)
svc_rec_score = recall_score(y_test, y_pred_test_svc)
svc_f1_score = f1_score(y_test, y_pred_test_svc)

print('Precision: %.3f' % svc_perc_score)
print('Recall: %.3f' % svc_rec_score)
print('F-measure: %.3f' % svc_f1_score)
```


Precision: 0.000
Recall: 0.000
F-measure: 0.000

```
In [74]: y_pred_prob_svc = svc.predict_proba(x_test)[:, 1]
svc_roc_auc_score = roc_auc_score(y_test, y_pred_prob_svc)
print('ROC AUC Score:', svc_roc_auc_score)
```

ROC AUC Score: 0.5

```
In [75]: skplt.metrics.plot_confusion_matrix(y_test, y_pred_test_svc, figsize=(6,6), cmap= 'BuPu')
```

```
Out[75]: <AxesSubplot:title={'center':'Confusion Matrix'}, xlabel='Predicted label', ylabel='True label'>
```



Random Forest Classification:

```
In [76]: rf = RandomForestClassifier(n_estimators = 100, criterion= 'entropy', random_state = 0)
rf.fit(x_data_balanced, y_data_balanced)
```

```
Out[76]: ▼ RandomForestClassifier
RandomForestClassifier(criterion='entropy', random_state=0)
```

```
In [77]: y_pred_train_rf = rf.predict(x_data_balanced)
acc_train_rf = accuracy_score(y_data_balanced, y_pred_train_rf)

y_pred_test_rf = rf.predict(x_test)
acc_test_rf = accuracy_score(y_test, y_pred_test_rf)
print(acc_train_rf)
print(acc_test_rf)
```

1.0
0.9383561643835616

```
In [78]: print(classification_report(y_test, y_pred_test_rf))
```

	precision	recall	f1-score	support
0	0.94	1.00	0.97	960
1	0.00	0.00	0.00	62

accuracy			0.94	1022
macro avg	0.47	0.50	0.48	1022
weighted avg	0.88	0.94	0.91	1022

```
In [79]: rf_perc_score = precision_score(y_test, y_pred_test_rf)
rf_rec_score= recall_score(y_test, y_pred_test_rf)
rf_f1_score = f1_score(y_test, y_pred_test_rf)

print('Precision: %.3f' %rf_perc_score )
print('Recall: %.3f' % rf_rec_score)
print('F-measure: %.3f' % rf_f1_score)
```


```
Precision: 0.000
Recall: 0.000
F-measure: 0.000
```

```
In [80]: y_pred_prob_rf = rf.predict_proba(x_test)[: , 1]
rf_roc_auc_score = roc_auc_score(y_test, y_pred_prob_rf)
print('ROC AUC Score:', rf_roc_auc_score)
```

```
ROC AUC Score: 0.6210853494623655
```

```
In [81]: skplt.metrics.plot_confusion_matrix(y_test, y_pred_test_rf, figsize=(6,6), cmap= 'Purple')
```

```
Out[81]: <AxesSubplot:title={'center':'Confusion Matrix'}, xlabel='Predicted label', ylabel='True label'>
```



GradientBoostingClassifier

```
In [82]: gbc = GradientBoostingClassifier()
gbc.fit(x_data_balanced,y_data_balanced)
```

```
Out[82]: ▾ GradientBoostingClassifier
GradientBoostingClassifier()
```

```
In [83]: y_pred_train_gbc = gbc.predict(x_data_balanced)
acc_train_gbc = accuracy_score(y_data_balanced, y_pred_train_gbc)
```

```

y_pred_test_gbc = gbc.predict(x_test)
acc_test_gbc = accuracy_score(y_test, y_pred_test_gbc)
print(acc_train_gbc)
print(acc_test_gbc)

```

0.9215384615384615

0.9393346379647749

In [84]: `print(classification_report(y_test, y_pred_test_gbc))`

```

              precision    recall  f1-score   support

     0           0.94         1.00         0.97         960
     1           0.00         0.00         0.00          62

 accuracy                   0.94         1022
 macro avg              0.47         0.50         0.48         1022
 weighted avg           0.88         0.94         0.91         1022

```

In [85]: `gbc_perc_score = precision_score(y_test, y_pred_test_gbc)`
`gbc_rec_score= recall_score(y_test, y_pred_test_gbc)`
`gbc_f1_score = f1_score(y_test, y_pred_test_gbc)`

```

print('Precision: %.3f' %gbc_perc_score )
print('Recall: %.3f' % gbc_rec_score)
print('F-measure: %.3f' % gbc_f1_score)

```

Precision: 0.000

Recall: 0.000


F-measure: 0.000

In [86]: `y_pred_prob_gbc = gbc.predict_proba(x_test)[: , 1]`
`gbc_roc_auc_score = roc_auc_score(y_test, y_pred_prob_gbc)`
`print('ROC AUC Score:', gbc_roc_auc_score)`

ROC AUC Score: 0.5873487903225807

In [87]: `skplt.metrics.plot_confusion_matrix(y_test, y_pred_test_gbc, figsize=(6,6), cmap= 'Purpl`

Out[87]: `<AxesSubplot:title={'center':'Confusion Matrix'}, xlabel='Predicted label', ylabel='True label'>`



XGB CLASSIFIER

```
In [88]: xgb = XGBClassifier(eval_metric= 'error', learning_rate= 0.05)
xgb.fit(x_data_balanced, y_data_balanced)
```

```
Out[88]: XGBClassifier
XGBClassifier(base_score=0.5, booster='gbtree', callbacks=None,
               colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1,
               early_stopping_rounds=None, enable_categorical=False,
               eval_metric='error', gamma=0, gpu_id=-1, grow_policy='depthwise',
               importance_type=None, interaction_constraints='',
               learning_rate=0.05, max_bin=256, max_cat_to_onehot=4,
               max_delta_step=0, max_depth=6, max_leaves=0, min_child_weight=1,
               missing=nan, monotone_constraints='()', n_estimators=100,
               n_jobs=0, num_parallel_tree=1, predictor='auto', random_state=0,
               reg_alpha=0, reg_lambda=1, ...)
```

```
In [89]: y_pred_train_xgb = xgb.predict(x_data_balanced)
acc_train_xgb = accuracy_score(y_data_balanced, y_pred_train_xgb)

y_pred_test_xgb = xgb.predict(x_test)
acc_test_xgb = accuracy_score(y_test, y_pred_test_xgb)

print(acc_train_xgb)
print(acc_test_xgb)

0.9443589743589743
0.9393346379647749
```

```
In [90]: print(classification_report(y_test, y_pred_test_xgb))
```

	precision	recall	f1-score	support
0	0.94	1.00	0.97	960
1	0.00	0.00	0.00	62
accuracy			0.94	1022
macro avg	0.47	0.50	0.48	1022
weighted avg	0.88	0.94	0.91	1022

```
In [91]: xgb_perc_score = precision_score(y_test, y_pred_test_xgb)
xgb_rec_score= recall_score(y_test, y_pred_test_xgb)
xgb_f1_score = f1_score(y_test, y_pred_test_xgb)

print('Precision: %.3f' %xgb_perc_score )
print('Recall: %.3f' % xgb_rec_score)
print('F-measure: %.3f' % xgb_f1_score)

Precision: 0.000
Recall: 0.000
F-measure: 0.000
```

```
In [93]: y_pred_prob_xgb = xgb.predict_proba(x_test)[: , 1]
xgb_roc_auc_score = roc_auc_score(y_test, y_pred_prob_xgb)
print('ROC AUC Score:', xgb_roc_auc_score)
```

ROC AUC Score: 0.6870547715053764

hyperparameters Tunning

```
In [94]: grid_models = [(LogisticRegression(), [{"C": np.logspace(-3, 3, 7), "penalty": ["l1", "l2"]}],
    (SVC(probability=True), [{"C": [10, 100], "gamma": [100, 500, 1000], 'kernel': ['linear', 'r
    (DecisionTreeClassifier(), [{"max_features": [5, 6, 10, 12, 14, 18, 20], 'max_depth': [6, 10, 12
    (RandomForestClassifier(), [{"n_estimators": [100, 150, 200], 'criterion': ['gini', 'entrop
    (XGBClassifier(), [{"learning_rate": [0.01, 0.05, 0.1, 0.5, 1], 'eval_metric': ['erro
```

```
In [95]: for i, j in grid_models:
    grid = GridSearchCV(estimator=i, param_grid = j, scoring = 'accuracy', cv = 5)
    grid.fit(x_data_balanced, y_data_balanced)
    best_accuracy = grid.best_score_
    best_param = grid.best_params_
    print('{}:\nBest Accuracy : {:.2f}%'.format(i, best_accuracy*100))
    print('Best Parameters : ', best_param)
    print('')
    print('-----')
    print('')
```

```
LogisticRegression():
Best Accuracy : 79.08%
Best Parameters : {'C': 100.0, 'penalty': 'l2'}
```

```
-----
```

```
SVC(probability=True):
Best Accuracy : 93.81%
Best Parameters : {'C': 10, 'gamma': 100, 'kernel': 'rbf', 'random_state': 0}
```

```
-----
```

```
DecisionTreeClassifier():
Best Accuracy : 91.59%
Best Parameters : {'criterion': 'entropy', 'max_depth': 20, 'max_features': 10, 'random
_state': 0}
```

```
-----
```

```
RandomForestClassifier():
Best Accuracy : 95.90%
Best Parameters : {'criterion': 'entropy', 'n_estimators': 150, 'random_state': 0}
```

```
-----
```

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
    colsample_bylevel=None, colsample_bynode=None,
    colsample_bytree=None, early_stopping_rounds=None,
    enable_categorical=False, eval_metric=None, gamma=None,
    gpu_id=None, grow_policy=None, importance_type=None,
    interaction_constraints=None, learning_rate=None, max_bin=None,
    max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
    max_leaves=None, min_child_weight=None, missing=nan,
    monotone_constraints=None, n_estimators=100, n_jobs=None,
    num_parallel_tree=None, predictor=None, random_state=None,
    reg_alpha=None, reg_lambda=None, ...):
Best Accuracy : 95.81%
Best Parameters : {'eval_metric': 'error', 'learning_rate': 0.5}
```

```
-----
```

```
In [96]: # classification with random forest

# Change parameters
rf = RandomForestClassifier(n_estimators = 100, criterion = 'entropy', random_state = 0)
rf.fit(x_data_balanced, y_data_balanced)
```

```
y_pred_train_rf = rf.predict(x_data_balanced)
acc_train_rf = accuracy_score(y_data_balanced, y_pred_train_rf)
```

```
y_pred_test_rf = rf.predict(x_test)
acc_test_rf = accuracy_score(y_test, y_pred_test_rf)
print(acc_train_rf)
print(acc_test_rf)
```

```
rf_perc_score = precision_score(y_test, y_pred_test_rf)
rf_rec_score = recall_score(y_test, y_pred_test_rf)
rf_f1_score = f1_score(y_test, y_pred_test_rf)
```


```
print('Precision: %.3f' % rf_perc_score)
print('Recall: %.3f' % rf_rec_score)
print('F-measure: %.3f' % rf_f1_score)
```

```
y_pred_prob_rf = rf.predict_proba(x_test)[:, 1]
rf_roc_auc_score = roc_auc_score(y_test, y_pred_prob_rf)
print('ROC AUC Score:', rf_roc_auc_score)
```

```
1.0
0.9383561643835616
Precision: 0.000
Recall: 0.000
F-measure: 0.000
ROC AUC Score: 0.6210853494623655
```

In [97]: `skplt.metrics.plot_confusion_matrix(y_test, y_pred_test_rf, figsize=(6,6), cmap= 'YlGnBu'`


Out[97]: `<AxesSubplot:title={'center':'Confusion Matrix'}, xlabel='Predicted label', ylabel='True label'>`



In [98]: `false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_pred_prob_rf)
roc_auc = auc(false_positive_rate, true_positive_rate)`

```
sns.set_theme(style = 'white')
plt.figure(figsize = (8, 8))
plt.plot(false_positive_rate, true_positive_rate, color = '#b01717', label = 'AUC = %0.3f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], linestyle = '--', color = '#174ab0')
plt.axis('tight')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
```

```
plt.legend()  
plt.show()
```

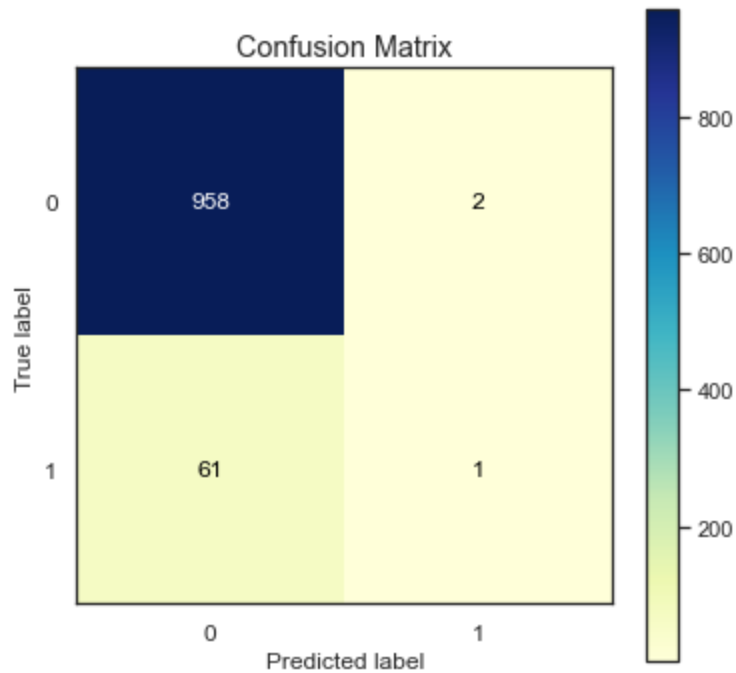


In [99]: *# Change parameters*

```
dt =DecisionTreeClassifier(max_features=5 , max_depth=12,criterion = 'entropy', random_s  
dt.fit(x_data_balanced, y_data_balanced)  
  
y_pred_train_dt = dt.predict(x_data_balanced)  
acc_train_dt = accuracy_score(y_data_balanced, y_pred_train_dt)  
  
y_pred_test_dt = dt.predict(x_test)  
acc_test_dt = accuracy_score(y_test, y_pred_test_dt)  
print(acc_train_dt)  
print(acc_test_dt)  
  
dt_perc_score = precision_score(y_test, y_pred_test_dt)  
dt_rec_score= recall_score(y_test, y_pred_test_dt)  
dt_f1_score = f1_score(y_test, y_pred_test_dt)  
  
print('Precision: %.3f' % dt_perc_score)  
print('Recall: %.3f' % dt_rec_score)  
print('F-measure: %.3f' % dt_f1_score)  
  
y_pred_prob_dt = dt.predict_proba(x_test)[: , 1]  
dt_roc_auc_score = roc_auc_score(y_test, y_pred_prob_dt)  
print('ROC AUC Score:', dt_roc_auc_score)
```

```
0.9235897435897436  
0.9383561643835616  
Precision: 0.333  
Recall: 0.016  
F-measure: 0.031  
ROC AUC Score: 0.5070228494623655
```


```
In [100... skplt.metrics.plot_confusion_matrix(y_test, y_pred_test_dt, figsize=(6,6), cmap= 'YlGnBu')
Out[100]: <AxesSubplot:title={'center':'Confusion Matrix'}, xlabel='Predicted label', ylabel='True label'>
```



```
In [101... # visualize Roc AUC Curve

false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_pred_prob_dt)
roc_auc = auc(false_positive_rate, true_positive_rate)

sns.set_theme(style = 'white')
plt.figure(figsize = (8, 8))
plt.plot(false_positive_rate,true_positive_rate, color = '#b01717', label = 'AUC = %0.3f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], linestyle = '--', color = '#174ab0')
plt.axis('tight')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.legend()
plt.show()
```



```
In [102... # Change parameters

xgb = XGBClassifier(eval_metric= 'error', learning_rate= 0.1, random_state=0)
xgb.fit(x_data_balanced, y_data_balanced)

y_pred_train_xgb = xgb.predict(x_data_balanced)
acc_train_xgb = accuracy_score(y_data_balanced, y_pred_train_xgb)

y_pred_test_xgb = xgb.predict(x_test)
acc_test_xgb = accuracy_score(y_test, y_pred_test_xgb)

print(acc_train_xgb)
print(acc_test_xgb)


xgb_perc_score = precision_score(y_test, y_pred_test_xgb)
xgb_rec_score= recall_score(y_test, y_pred_test_xgb)
xgb_f1_score = f1_score(y_test, y_pred_test_xgb)

print('Precision: %.3f' %xgb_perc_score )
print('Recall: %.3f' % xgb_rec_score)
print('F-measure: %.3f' % xgb_f1_score)

y_pred_prob_xgb = xgb.predict_proba(x_test)[:, 1]
xgb_roc_auc_score = roc_auc_score(y_test, y_pred_prob_xgb)
print('ROC AUC Score:', xgb_roc_auc_score)

0.973974358974359
0.9393346379647749
Precision: 0.000
Recall: 0.000
F-measure: 0.000
ROC AUC Score: 0.6698840725806452
```


```
In [103... skplt.metrics.plot_confusion_matrix(y_test, y_pred_test_xgb, figsize=(6,6), cmap= 'YlGnB
```



```
In [104... # visualize Roc AUC Curve

false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_pred_prob_xgb)
roc_auc = auc(false_positive_rate, true_positive_rate)

sns.set_theme(style = 'white')
plt.figure(figsize = (8, 8))
plt.plot(false_positive_rate,true_positive_rate, color = '#b01717', label = 'AUC = %0.3f' % roc_auc)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], linestyle = '--', color = '#174ab0')
plt.axis('tight')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.legend()
plt.show()
```



```
In [ ]: svc = SVC(C=10, gamma=1000 ,probability=True)
svc.fit(x_data_balanced, y_data_balanced)

y_pred_train_svc = svc.predict(x_data_balanced)
acc_train_svc = accuracy_score(y_data_balanced, y_pred_train_svc)

y_pred_test_svc = svc.predict(x_test)
acc_test_svc = accuracy_score(y_test, y_pred_test_svc)

print(acc_train_svc)
print(acc_test_svc)

svc_perc_score = precision_score(y_test, y_pred_test_svc)
svc_rec_score= recall_score(y_test, y_pred_test_svc)
svc_f1_score = f1_score(y_test, y_pred_test_svc)

print('Precision: %.3f' % svc_perc_score)
print('Recall: %.3f' % svc_rec_score)
print('F-measure: %.3f' % svc_f1_score)

y_pred_prob_svc = svc.predict_proba(x_test)[: , 1]
svc_roc_auc_score= roc_auc_score(y_test, y_pred_prob_svc)
print('ROC AUC Score:', svc_roc_auc_score)
```

```
In [ ]: skplt.metrics.plot_confusion_matrix(y_test, y_pred_test_svc, figsize=(6,6), cmap= 'YlGnB
```

```
In [ ]: # visualize Roc AUC Curve

false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_pred_prob_svc)
roc_auc = auc(false_positive_rate, true_positive_rate)

sns.set_theme(style = 'white')
plt.figure(figsize = (8, 8))
```



```
plt.plot(false_positive_rate,true_positive_rate, color = '#b01717', label = 'AUC = %0.3f')
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], linestyle = '--', color = '#174ab0')
plt.axis('tight')
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.legend()
plt.show()
```

In []: